

Adopting Karpenter for Cost and Simplicity



Logan Ballard
Senior Software Engineer

Node Summary

Node Name	Pods	CPU Utilization	Memory Utilization	
	25	<div><div></div></div> 94.8%	<div><div></div></div> 95.9%	C
	35	<div><div></div></div> 97.3%	<div><div></div></div> 90.5%	o
	24	<div><div></div></div> 92.3%	<div><div></div></div> 99.7%	o
	21	<div><div></div></div> 98.6%	<div><div></div></div> 55.9%	o
	30	<div><div></div></div> 97.6%	<div><div></div></div> 90.3%	o
	35	<div><div></div></div> 98.9%	<div><div></div></div> 100.0%	o
	33	<div><div></div></div> 98.8%	<div><div></div></div> 93.0%	o
	34	<div><div></div></div> 99.7%	<div><div></div></div> 97.4%	o
	25	<div><div></div></div> 97.6%	<div><div></div></div> 63.9%	o
	36	<div><div></div></div> 97.3%	<div><div></div></div> 91.6%	o
	31	<div><div></div></div> 99.6%	<div><div></div></div> 98.8%	o
	30	<div><div></div></div> 98.9%	<div><div></div></div> 98.6%	o
	37	<div><div></div></div> 98.1%	<div><div></div></div> 96.1%	o
	32	<div><div></div></div> 98.7%	<div><div></div></div> 96.8%	o
	32	<div><div></div></div> 96.1%	<div><div></div></div> 97.8%	o
	31	<div><div></div></div> 99.0%	<div><div></div></div> 98.9%	o
	33	<div><div></div></div> 96.7%	<div><div></div></div> 98.1%	o
	37	<div><div></div></div> 99.5%	<div><div></div></div> 96.1%	o
	32	<div><div></div></div> 98.3%	<div><div></div></div> 96.4%	o



Logan Ballard

Senior Software Engineer



Adopting Karpenter for Cost and Simplicity

Optimizing Autoscaling

When Cluster Autoscaler isn't meeting your needs in EKS, consider Karpenter. We will walk through our journey to Karpenter, alternatives we considered, associated trade offs, and why Karpenter ultimately has helped us reduce costs and complexity to better serve our customers on AWS.



Overview of platform engineering at Grafana Labs

- Company goal: one-stop-shop for observability
 - “Query, visualize, alert on, and understand your data no matter where it’s stored. With Grafana you can create, explore, and share all of your data through beautiful, flexible dashboards.”



Overview of platform engineering at Grafana Labs

- Company goal: one-stop-shop for observability
 - “Query, visualize, alert on, and understand your data no matter where it’s stored. With Grafana you can create, explore, and share all of your data through beautiful, flexible dashboards.”
- Grafana platform team provides a PaaS for the product layer
 - Managing k8s, CI/CD, and interfaces for the “golden path”



Overview of platform engineering at Grafana Labs

- Company goal: one-stop-shop for observability
 - “Query, visualize, alert on, and understand your data no matter where it’s stored. With Grafana you can create, explore, and share all of your data through beautiful, flexible dashboards.”
- Grafana platform team provides a PaaS for the product layer
 - Managing k8s, CI/CD, and interfaces for the “golden path”
- To deliver the platform, we use managed Kubernetes offerings
 - GKE (Google)
 - EKS (Amazon)
 - AKS (Azure)



Overview of platform engineering at Grafana Labs

- Company goal: one-stop-shop for observability
 - “Query, visualize, alert on, and understand your data no matter where it’s stored. With Grafana you can create, explore, and share all of your data through beautiful, flexible dashboards.”
- Grafana platform team provides a PaaS for the product layer
 - Managing k8s, CI/CD, and interfaces for the “golden path”
- To deliver the platform, we use managed Kubernetes offerings
 - GKE (Google)
 - EKS (Amazon)
 - AKS (Azure)
- Platform team is further divided into sub-specialties, one of which is capacity



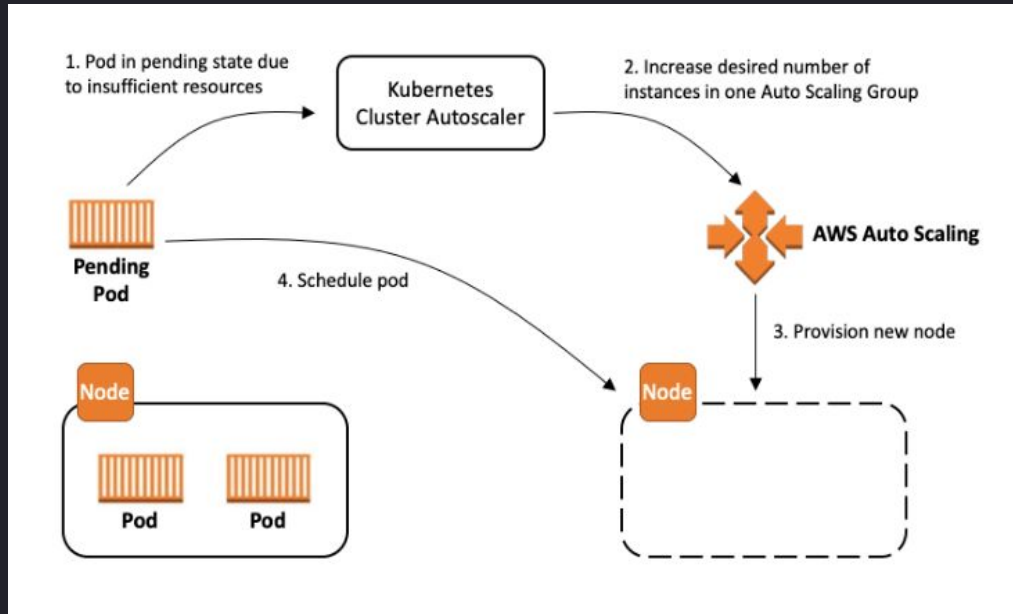
Overview of capacity at Grafana Labs

- Capacity is responsible for engineering around usage optimization as well as reliability
 - Observability into cost and utilization
 - Ownership of scaling tools
 - Company-wide efforts to optimize resource usage
- Today we'll be focusing on autoscaling



Overview of Cluster Autoscaler

- Up until last year, we used Cluster Autoscaler exclusively
- Scaling heuristic:
 - Demand goes up and down in the form of pod resource requests
 - If there is too much CPU or Memory requested, more nodes are added
 - If resource requests scale down, nodes are removed



Source: <https://aws.github.io/aws-eks-best-practices/cluster-autoscaling/>



What we liked about Cluster Autoscaler

- Simple, standardized, and has lots of support



What we liked about Cluster Autoscaler

- Simple, standardized, and has lots of support
- Accomplished basic scaling needs



What we liked about Cluster Autoscaler

- Simple, standardized, and has lots of support
- Accomplished basic scaling needs
- Infrastructure was managed in a standardized way
 - Node Groups defined in Terraform



What we liked about Cluster Autoscaler

- Simple, standardized, and has lots of support
- Accomplished basic scaling needs
- Infrastructure was managed in a standardized way
 - Node Groups defined in Terraform
- Deployment was managed in a standardized way, as either:
 - GKE: part of the managed offering
 - EKS: as-code in jsonnet, same as all our other workloads



What we liked about Cluster Autoscaler

- Simple, standardized, and has lots of support
- Accomplished basic scaling needs
- Infrastructure was managed in a standardized way
 - Node Groups defined in Terraform
- Deployment was managed in a standardized way, as either:
 - GKE: part of the managed offering
 - EKS: as-code in jsonnet, same as all our other workloads
- ...we already had it



What we didn't like about Cluster Autoscaler

- Lots of overhead when defining diverse machine types



What we didn't like about Cluster Autoscaler

- Lots of overhead when defining diverse machine types
- Inability to understand max pod IP restrictions when making scaling decisions



What we didn't like about Cluster Autoscaler

- Lots of overhead when defining diverse machine types
- Inability to understand max pod IP restrictions when making scaling decisions
- No support for “soft” scheduling constraints
 - E.g. `preferredDuringSchedulingIgnoredDuringExecution`



What we didn't like about Cluster Autoscaler

- Lots of overhead when defining diverse machine types
- Inability to understand max pod IP restrictions when making scaling decisions
- No support for “soft” scheduling constraints
 - E.g. `preferredDuringSchedulingIgnoredDuringExecution`
- Experimentation was challenging



What we didn't like about Cluster Autoscaler

- Lots of overhead when defining diverse machine types
- Inability to understand max pod IP restrictions when making scaling decisions
- No support for “soft” scheduling constraints
 - E.g. `preferredDuringSchedulingIgnoredDuringExecution`
- Experimentation was challenging
- Upgrades were a long process



What we didn't like about Cluster Autoscaler

- Lots of overhead when defining diverse machine types
- Inability to understand max pod IP restrictions when making scaling decisions
- No support for “soft” scheduling constraints
 - E.g. `preferredDuringSchedulingIgnoredDuringExecution`
- Experimentation was challenging
- Upgrades were a long process
- No fallback for workloads from spot nodes -> on demand nodes



Why change?



Why change?

- We measure “idle ratio” to calculate scheduling efficiency



Why change?

- We measure “idle ratio” to calculate scheduling efficiency
- Idle ratio is calculated as a product of:
 - $(\text{CPU allocatable}) - (\text{CPU Used}) / (\text{CPU allocatable})$
 - $(\text{Memory allocatable}) - (\text{Memory Used}) / (\text{Memory allocatable})$



Why change?

- We measure “idle ratio” to calculate scheduling efficiency
- Idle ratio is calculated as a product of:
 - $(\text{CPU allocatable}) - (\text{CPU Used}) / (\text{CPU allocatable})$
 - $(\text{Memory allocatable}) - (\text{Memory Used}) / (\text{Memory allocatable})$
- Idle ratio is a “golden signal” for capacity squad
- High idle ratio is wasted \$\$\$



Why change?

- We measure “idle ratio” to calculate scheduling efficiency
- Idle ratio is calculated as a product of:
 - $(\text{CPU allocatable}) - (\text{CPU Used}) / (\text{CPU allocatable})$
 - $(\text{Memory allocatable}) - (\text{Memory Used}) / (\text{Memory allocatable})$
- Idle ratio is a “golden signal” for capacity squad
- High idle ratio is wasted \$\$\$




Idle Ratio Visualized

GKE

vs

AWS





Idle Ratio Visualized

GKE

vs

AWS





Idle Ratio Visualized

GKE

vs

AWS



Idle ratio in AWS was almost double that of GKE



What to do?



What to do?

- Nothing?



What to do?

- Nothing?
- Try to replicate the setup of GKE with it's OPTIMIZE_UTILIZATION autoscaling profile
 - Run a custom scheduler: 'MostAllocated' scoring strategy
 - Configure Cluster Autoscaler: set expanders to priority? Most pods?
 - All of this is obscured behind the GKE control plane and would be a tedious process



What to do?

- Nothing?
- Try to replicate the setup of GKE with it's OPTIMIZE_UTILIZATION autoscaling profile
 - Run a custom scheduler: 'MostAllocated' scoring strategy
 - Configure Cluster Autoscaler: set expanders to priority? Most pods?
 - All of this is obscured behind the GKE control plane and would be a tedious process
- Manage our own control plane?



What to do?

- Nothing?
- Try to replicate the setup of GKE with it's OPTIMIZE_UTILIZATION autoscaling profile
 - Run a custom scheduler: 'MostAllocated' scoring strategy
 - Configure Cluster Autoscaler: set expanders to priority? Most pods?
 - All of this is obscured behind the GKE control plane and would be a tedious process
- Manage our own control plane?
 - No



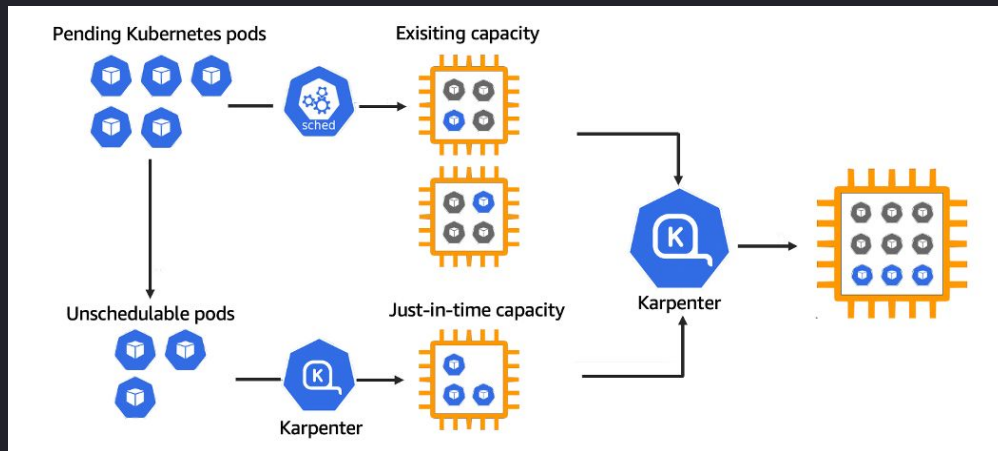
What to do?

- Nothing?
- Try to replicate the setup of GKE with it's OPTIMIZE_UTILIZATION autoscaling profile
 - Run a custom scheduler: 'MostAllocated' scoring strategy
 - Configure Cluster Autoscaler: set expanders to priority? Most pods?
 - All of this is obscured behind the GKE control plane and would be a tedious process
- Manage our own control plane?
 - No
- Karpenter?



Overview of Karpenter

- Open Source
- EKS specific
 - Update: Core logic of Karpenter is now part of the autoscaling SIG, as of November 2023
- “Just in time” model for upscaling
 - Matches workload size and shape to machines
- “Consolidation” model for downscaling
 - Continual bin-packing
- High level of visibility into decisions
- ...we were sold



Source:

<https://aws.amazon.com/blogs/aws/introducing-karpenter-an-open-source-high-performance-kubernetes-cluster-autoscaler/>



Goals to get there



Goals to get there

- Zero or minimal downtime



Goals to get there

- Zero or minimal downtime
- Transparent to end-users of platform



Goals to get there

- Zero or minimal downtime
- Transparent to end-users of platform
- Reproducible migration across different environments



Goals to get there

- Zero or minimal downtime
- Transparent to end-users of platform
- Reproducible migration across different environments
- Observability in all steps of the migration



Migration details



Migration details

- Define all resources in code
 - Supporting Infrastructure (IAM, SQS) defined as Terraform
 - Kubernetes resources (deployment, service account, etc) defined in jsonnet



Migration details

- Define all resources in code
 - Supporting Infrastructure (IAM, SQS) defined as Terraform
 - Kubernetes resources (deployment, service account, etc) defined in jsonnet
- Migration ordering of events is very important



Migration details

- Define all resources in code
 - Supporting Infrastructure (IAM, SQS) defined as Terraform
 - Kubernetes resources (deployment, service account, etc) defined in jsonnet
- Migration ordering of events is very important
 - Create “critical” node group for Karpenter itself to run on



Migration details

- Define all resources in code
 - Supporting Infrastructure (IAM, SQS) defined as Terraform
 - Kubernetes resources (deployment, service account, etc) defined in jsonnet
- Migration ordering of events is very important
 - Create “critical” node group for Karpenter itself to run on
 - Deploy all Karpenter resources except provisioners and node templates



Migration details

- Define all resources in code
 - Supporting Infrastructure (IAM, SQS) defined as Terraform
 - Kubernetes resources (deployment, service account, etc) defined in jsonnet
- Migration ordering of events is very important
 - Create “critical” node group for Karpenter itself to run on
 - Deploy all Karpenter resources except provisioners and node templates
 - Turn off Cluster Autoscaler
 - Deploy Karpenter Provisioners + Node Templates



Migration details

- Define all resources in code
 - Supporting Infrastructure (IAM, SQS) defined as Terraform
 - Kubernetes resources (deployment, service account, etc) defined in jsonnet
- Migration ordering of events is very important
 - Create “critical” node group for Karpenter itself to run on
 - Deploy all Karpenter resources except provisioners and node templates
 - Turn off Cluster Autoscaler
 - Deploy Karpenter Provisioners + Node Templates
 - Gradually cordon and drain existing nodes



Migration details

- Define all resources in code
 - Supporting Infrastructure (IAM, SQS) defined as Terraform
 - Kubernetes resources (deployment, service account, etc) defined in jsonnet
- Migration ordering of events is very important
 - Create “critical” node group for Karpenter itself to run on
 - Deploy all Karpenter resources except provisioners and node templates
 - Turn off Cluster Autoscaler
 - Deploy Karpenter Provisioners + Node Templates
 - Gradually cordon and drain existing nodes
 - Delete existing node groups



Results



Results

- Migration accomplished without downtime



Results

- Migration accomplished without downtime
- Massively lowered idle ratio, costs
 - More effective distribution of resources
 - Allowed us to take advantage of varying machine types



Results

- Migration accomplished without downtime
- Massively lowered idle ratio, costs
 - More effective distribution of resources
 - Allowed us to take advantage of varying machine types
- Improved reliability



Results

- Migration accomplished without downtime
- Massively lowered idle ratio, costs
 - More effective distribution of resources
 - Allowed us to take advantage of varying machine types
- Improved reliability
- Better developer experience




Results

- Migration accomplished without downtime
- Massively lowered idle ratio, costs
 - More effective distribution of resources
 - Allowed us to take advantage of varying machine types
- Improved reliability
- Better developer experience
- Simplified infrastructure



Results - idle ratio visualized

- Costs lowered
- Idle ratio lowered



Results: cost efficiency visualized



Large cluster (~350 nodes): pre-Karpenter



Results: cost efficiency visualized



Large cluster (~350 nodes): post-Karpenter



Results - improved reliability

Before

Rigid Options

- Workloads could either require spot or on-demand
- Maxing out autoscaling groups would result in alerts, downtime
- Single instance node groups subject to scarcity

vs

After

More Flexibility

- Workloads can *prefer* spot, but fall back to on-demand
- Fleet is not bounded by autoscaling group limitations
- Wide variety of instance types all but guarantees availability



Results - developer experience

Before

High Barrier to Experimentation

- Trying out new instance types required lots of overhead
- Infrastructure managed in Terraform
- Specialized workloads required specialized node groups

vs

After

Easy to Play Around

- New instance types are easy to try out
- Infrastructure managed by Kubernetes object
- Expressive “requirements” syntax for Karpenter allows workloads to ask for what they want, not peg themselves to a machine type



Results - simplified infrastructure

Before

Slow and Manual

- Upgrades handled on a node group basis
- Definition of cluster in Terraform was largely copy/paste

vs

After

Fast and Automated

- Upgrades can be hands-off
- Definition of cluster in Karpenter is largely programmatic



Trade-offs



Trade-offs

- No unified autoscaling solution across cloud providers, EKS only
 - Although... Azure is under development!



Trade-offs

- No unified autoscaling solution across cloud providers, EKS only
 - Although... Azure is under development!
- Karpenter is still considered “beta”
 - Recent release has a breaking API change



Trade-offs

- No unified autoscaling solution across cloud providers, EKS only
 - Although... Azure is under development!
- Karpenter is still considered “beta”
 - Recent release has a breaking API change
- No spot-to-spot consolidation at the moment



Trade-offs

- No unified autoscaling solution across cloud providers, EKS only
 - Although... Azure is under development!
- Karpenter is still considered “beta”
 - Recent release has a breaking API change
- No spot-to-spot consolidation at the moment
- Deprovisioning lacks fine-grained control



Wrap-up



Wrap-up

- Measuring idleness helps greatly for cost optimization



Wrap-up

- Measuring idleness helps greatly for cost optimization
- Karpenter requires buy-in, but is a significant EKS opportunity



Wrap-up

- Measuring idleness helps greatly for cost optimization
- Karpenter requires buy-in, but is a significant EKS opportunity
- Ensuring a smooth migration was high effort, but worth it



Wrap-up

- Measuring idleness helps greatly for cost optimization
- Karpenter requires buy-in, but is a significant EKS opportunity
- Ensuring a smooth migration was high effort, but worth it
- Passing information on to end-users of the platform yielded and continues to yield high benefits



Wrap-up

- Measuring idleness helps greatly for cost optimization
- Karpenter requires buy-in, but is a significant EKS opportunity
- Ensuring a smooth migration was high effort, but worth it
- Passing information on to end-users of the platform yielded and continues to yield high benefits
- We are excited to see where this tool goes!



Questions?



Additional Info

- Contact me
 - loganballard@gmail.com
 - <https://github.com/logyball>
- Karpenter
 - <https://karpenter.sh/>
- Grafana
 - *"How Grafana Labs switched to Karpenter to reduce costs and complexities in Amazon EKS"*
 - <https://bit.ly/grafana-karpenter>



QR code to relevant blog post





Take Grafana Labs
Observability Survey